

# Neural Networks & Deep Learning - ICP-9

CS 5720 (CRN 23216)

Student ID: 700745451

Student Name: Kamala Ramesh

Execution of the provided Autoencoder

```
✓ [1] from keras.layers import Input, Dense
11s   from keras.models import Model
      from keras.datasets import mnist, fashion_mnist
      import numpy as np

✓ [2] # this is the size of our encoded representations
0s    encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

      # this is our input placeholder
      input_img = Input(shape=(784,))
      # "encoded" is the encoded representation of the input
      encoded = Dense(encoding_dim, activation='relu')(input_img)

      # "decoded" is the lossy reconstruction of the input
      decoded = Dense(784, activation='sigmoid')(encoded)

✓ [3] # this model maps an input to its reconstruction
0s    autoencoder = Model(input_img, decoded)
      # this model maps an input to its encoded representation
      autoencoder.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

✓ [4] (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
2s    x_train = x_train.astype('float32') / 255.
      x_test = x_test.astype('float32') / 255.
      x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
      x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
      29515/29515 [=====] - 0s 0us/step
      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
      26421880/26421880 [=====] - 1s 0us/step
      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
      5148/5148 [=====] - 0s 0us/step
      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
      4422102/4422102 [=====] - 0s 0us/step
```

```

21s ▶ autoencoder.fit(x_train, x_train,
                    epochs=5,
                    batch_size=256,
                    shuffle=True,
                    validation_data=(x_test, x_test))

Epoch 1/5
235/235 [=====] - 7s 19ms/step - loss: 0.4139 - accuracy: 0.0086 - val_loss: 0.3431 - val_accuracy: 0.0127
Epoch 2/5
235/235 [=====] - 4s 18ms/step - loss: 0.3285 - accuracy: 0.0138 - val_loss: 0.3206 - val_accuracy: 0.0128
Epoch 3/5
235/235 [=====] - 3s 13ms/step - loss: 0.3120 - accuracy: 0.0137 - val_loss: 0.3083 - val_accuracy: 0.0148
Epoch 4/5
235/235 [=====] - 3s 13ms/step - loss: 0.3020 - accuracy: 0.0162 - val_loss: 0.3002 - val_accuracy: 0.0164
Epoch 5/5
235/235 [=====] - 3s 15ms/step - loss: 0.2954 - accuracy: 0.0184 - val_loss: 0.2953 - val_accuracy: 0.0185
<keras.callbacks.History at 0x7faa4e3f0a30>

```

## 1. To add one hidden layer to the existing code

```

5s [1] from keras.layers import Input, Dense
      from keras.models import Model
      from keras.datasets import mnist, fashion_mnist
      import numpy as np
      import matplotlib.pyplot as plt

```

```

0s [2] # this is the size of our encoded representations
      encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

      # this is our input placeholder
      input_img = Input(shape=(784,))
      # "encoded" is the encoded representation of the input
      encoded = Dense(encoding_dim, activation='relu')(input_img)

      # Define the hidden layer
      hidden = Dense(784, activation='relu')(encoded)

      # "decoded" is the lossy reconstruction of the input
      decoded = Dense(784, activation='sigmoid')(hidden)

```

```

0s [3] # this model maps an input to its reconstruction
      autoencoder = Model(input_img, decoded)
      # this model maps an input to its encoded representation
      autoencoder.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

```

```

1s ▶ (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
      x_train = x_train.astype('float32') / 255.
      x_test = x_test.astype('float32') / 255.
      x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
      x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
      29515/29515 [=====] - 0s 0us/step
      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
      26421880/26421880 [=====] - 0s 0us/step
      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
      5148/5148 [=====] - 0s 0us/step
      Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
      4422102/4422102 [=====] - 0s 0us/step

```

```

✓ 1m [5] history = autoencoder.fit(x_train, x_train,
                                epochs=5,
                                batch_size=256,
                                shuffle=True,
                                validation_data=(x_test, x_test))

Epoch 1/5
235/235 [=====] - 11s 42ms/step - loss: 0.3535 - accuracy: 0.0125 - val_loss: 0.3054 - val_accuracy: 0.0168
Epoch 2/5
235/235 [=====] - 9s 37ms/step - loss: 0.2957 - accuracy: 0.0155 - val_loss: 0.2923 - val_accuracy: 0.0201
Epoch 3/5
235/235 [=====] - 8s 35ms/step - loss: 0.2870 - accuracy: 0.0200 - val_loss: 0.2867 - val_accuracy: 0.0227
Epoch 4/5
235/235 [=====] - 9s 39ms/step - loss: 0.2826 - accuracy: 0.0236 - val_loss: 0.2833 - val_accuracy: 0.0235
Epoch 5/5
235/235 [=====] - 10s 41ms/step - loss: 0.2798 - accuracy: 0.0274 - val_loss: 0.2808 - val_accuracy: 0.0290

```

## 2. Prediction on reconstructed data and visualize the reconstructed data


```

✓ 2s [6] reconstructed_data = autoencoder.predict(x_test)


313/313 [=====] - 2s 5ms/step

```

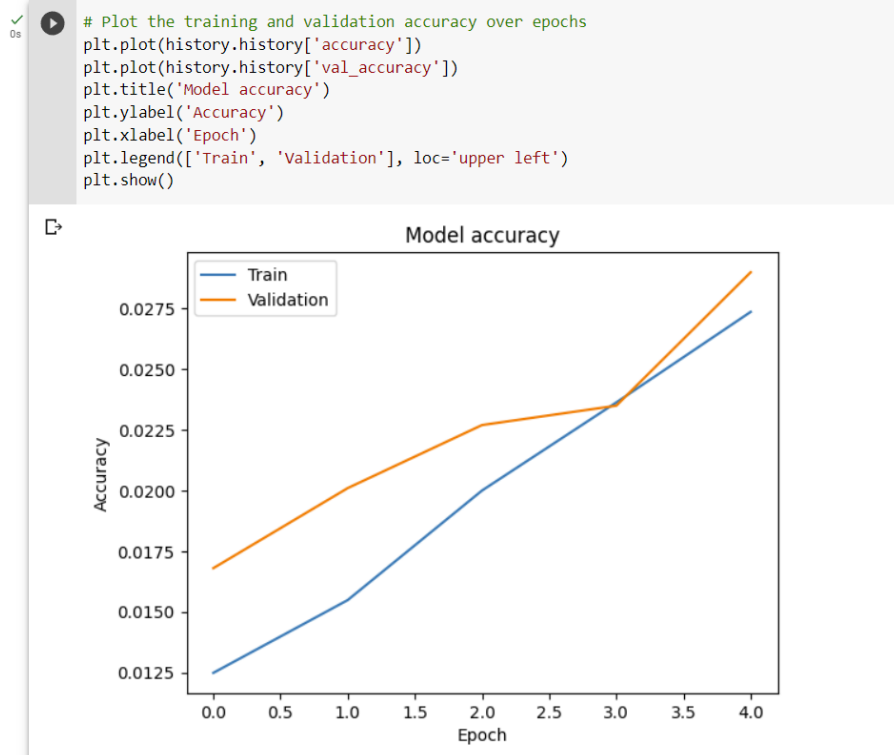
Visualize the original data

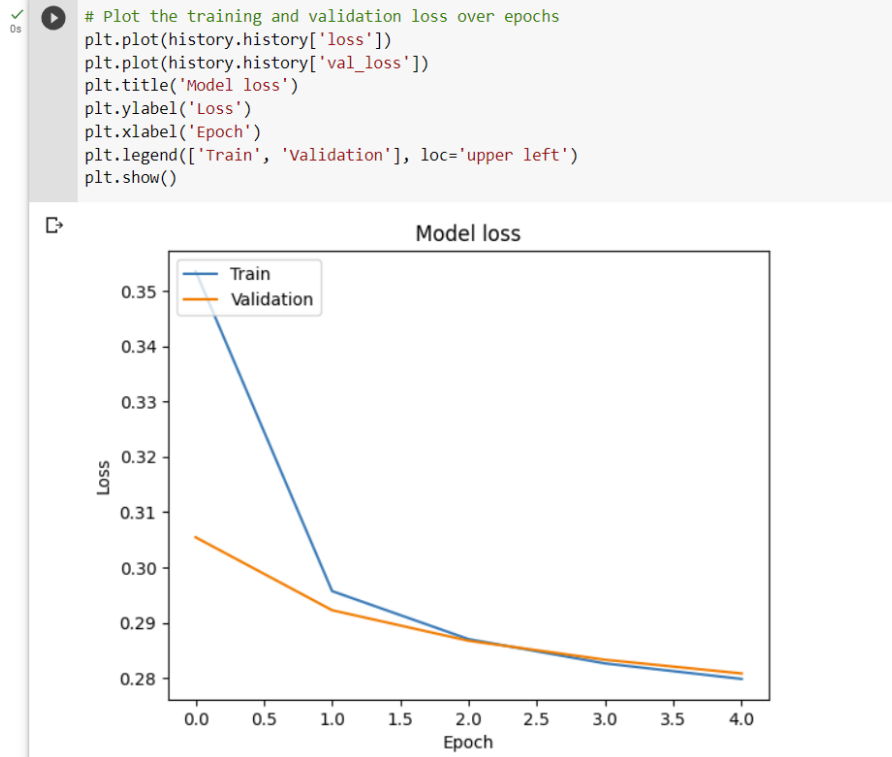


Visualize the reconstructed data



Accuracy and Loss Plot





### 3. Repeating the above scenarios for Denoising Encoder

```

6s ✓ [1] from keras.layers import Input, Dense
    from keras.models import Model
    from keras.datasets import fashion_mnist
    import numpy as np
    import matplotlib.pyplot as plt

```

```

0s ✓ [2] # this is the size of our encoded representations
    encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

    # this is our input placeholder
    input_img = Input(shape=(784,))
    # "encoded" is the encoded representation of the input
    encoded = Dense(encoding_dim, activation='relu')(input_img)

    # Define the hidden layer
    hidden = Dense(784, activation='relu')(encoded)

    # "decoded" is the lossy reconstruction of the input
    decoded = Dense(784, activation='sigmoid')(hidden)

```

```

0s ✓ [3] # this model maps an input to its reconstruction
    autoencoder = Model(input_img, decoded)
    # this model maps an input to its encoded representation
    autoencoder.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

```

✓  
9s

```
(x_train, _), (x_test, _) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

#introducing noise
noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)
```

📁 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>  
29515/29515 [=====] - 0s 0us/step  
Download data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>  
26421880/26421880 [=====] - 1s 0us/step  
Download data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>  
5148/5148 [=====] - 0s 0us/step  
Download data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>  
4422102/4422102 [=====] - 0s 0us/step

✓  
1m

```
history = autoencoder.fit(x_train_noisy, x_train,
                          epochs=10,
                          batch_size=128,
                          shuffle=True,
                          validation_data=(x_test_noisy, x_test_noisy))
```


📁 Epoch 1/10  
469/469 [=====] - 13s 25ms/step - loss: 0.3414 - accuracy: 0.0120 - val\_loss: 0.2856 - val\_accuracy: 0.0088  
Epoch 2/10  
469/469 [=====] - 11s 24ms/step - loss: 0.3041 - accuracy: 0.0150 - val\_loss: 0.2660 - val\_accuracy: 0.0101  
Epoch 3/10  
469/469 [=====] - 12s 25ms/step - loss: 0.2975 - accuracy: 0.0176 - val\_loss: 0.2551 - val\_accuracy: 0.0106  
Epoch 4/10  
469/469 [=====] - 11s 23ms/step - loss: 0.2940 - accuracy: 0.0203 - val\_loss: 0.2479 - val\_accuracy: 0.0099  
Epoch 5/10  
469/469 [=====] - 11s 23ms/step - loss: 0.2920 - accuracy: 0.0231 - val\_loss: 0.2436 - val\_accuracy: 0.0129  
Epoch 6/10  
469/469 [=====] - 12s 26ms/step - loss: 0.2907 - accuracy: 0.0242 - val\_loss: 0.2405 - val\_accuracy: 0.0118  
Epoch 7/10  
469/469 [=====] - 12s 26ms/step - loss: 0.2899 - accuracy: 0.0259 - val\_loss: 0.2386 - val\_accuracy: 0.0113  
Epoch 8/10  
469/469 [=====] - 12s 25ms/step - loss: 0.2894 - accuracy: 0.0277 - val\_loss: 0.2376 - val\_accuracy: 0.0114  
Epoch 9/10  
469/469 [=====] - 12s 25ms/step - loss: 0.2889 - accuracy: 0.0277 - val\_loss: 0.2358 - val\_accuracy: 0.0112  
Epoch 10/10  
469/469 [=====] - 12s 25ms/step - loss: 0.2885 - accuracy: 0.0280 - val\_loss: 0.2351 - val\_accuracy: 0.0112

✓  
1s


```
reconstructed_data = autoencoder.predict(x_test_noisy)
```

313/313 [=====] - 1s 3ms/step

Visualize the original data




Visualize reconstructed data



## Accuracy and Loss plots

```
✓ 1s # Plot the training and validation accuracy over epochs
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```





✓  
0s

```
# Plot the training and validation loss over epochs
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```

